

Exercise – Using the Spark DataFrame API

Objective

The objectives for this exercise are:

1. Gain familiarity with creating a SparkSession with pyspark
2. Use the Spark DataFrame API to query a dataset
3. Gain familiarity with converting a Spark DataFrame to a Pandas DataFrame

Steps

1. Create an empty Jupyter notebook
2. Create a SparkSession object with master set to "local[*]"
3. Read the csv file "countries.csv" into a Spark DataFrame
4. Use the Spark DataFrame API to answer the following questions about the data that was just loaded:
 - a. Display the first 15 rows
 - b. Display all rows that have the region "Antarctica" (hint: use the 'filter' function)
 - c. How many unique entries are in the Code field? (hint: use the 'select' function)
 - d. Create a tempView so that you can query the data with SQL statements
 - e. Using the tempView display the unique values of the "GovernmentForm" column
 - f. Convert the data type of the LifeExpectancy, Population & GNP columns to the type "pyspark.sql.types.DoubleType"
 - g. What is the average population per continent (hint: use the "groupby" and "avg" functions)
 - h. Convert the Spark DataFrame to a pandas DataFrame and display the first 5 rows
 - i. Optional: use the pandas DataFrame to create a line chart of LifeExpectancy Vs GNP